

EC 440 – Introduction to Operating Systems

Manuel Egele

Department of Electrical &
Computer Engineering
Boston University

Process vs. Thread

Address Space (Data/Heap)

i = 42

P1

i = 23

P2

Address Space (Stack)

17

11

P1

17

11

P2

Registers

PC = 5

P1

PC = 6

P2

Running

Address Space (Data/Heap)

i = 23

Address Space (Stack)

17

T1

T2

Registers

PC = 10

T1

PC = 16

T2

Running

Thread Primitives (e.g., pthread_create)

\$ man pthread_create

PTHREAD_CREATE(3)

Linux Programmer's Manual

NAME

pthread_create - create a new thread

SYNOPSIS

```
#include <pthread.h>
```

```
int pthread_create(pthread_t *thread, const pthread_attr_t *attr,  
                  void *(*start_routine) (void *), void *arg);
```

Compile and link with -pthread.

Each thread has a separate stack!

The Stack ...

Procedures

```
int f(int x) {return    x + 1; }  
int g(int x) {return    f(x); }  
int h(int x) {return f(x *2); }
```

Procedures (functions) are intrinsically linked to the stack

- Provides space for local variables
- Records where to return to
- Used to pass arguments (sometimes)


Implemented using stack frames

- Also known as activation records

Procedures: Calling and Returning


Instruction	Effect	Description
call x	$\text{rsp} \leftarrow \text{rsp} - 8$	Decrement rsp by 8
	$\text{Mem}(\text{rsp}) \leftarrow \text{Succ}(\text{rip})$	Store successor
	$\text{rip} \leftarrow \text{Addr}(x)$	Jump to address
ret	$\text{rip} \leftarrow \text{Mem}(\text{rsp})$	Pop successor into rip
	$\text{rsp} \leftarrow \text{rsp} + 8$	Increment rsp by 8

Stack Frame




```
1  int auth(const char* user) {  
2      size_t i;  
3      char buf[16];  
4      strncpy(buf, user, sizeof(buf));  
5      buf[sizeof(buf) - 1] = '\\0';  
6      for (i = 0; i < sizeof(buf); i++)  
7          buf[i] ^= 0xe5;  
8      return !memcmp(buf, "secret", 6);  
9  }
```

Stack Frame



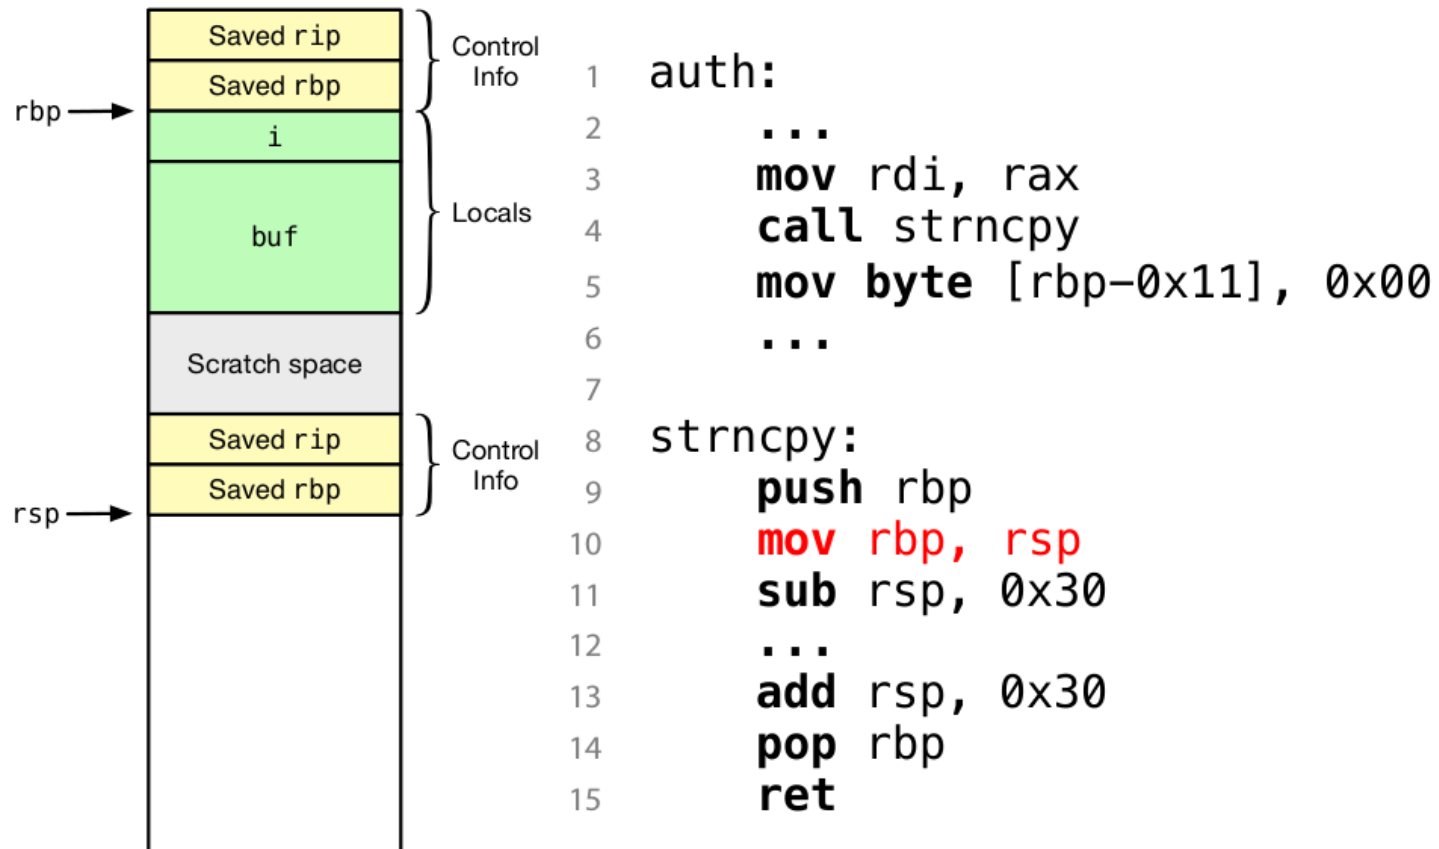
```
1  auth:
2      ...
3      mov rdi, rax
4      call strncpy
5      mov byte [rbp-0x11], 0x00
6      ...
7
8  strncpy:
9      push rbp
10     mov rbp, rsp
11     sub rsp, 0x30
12     ...
13     add rsp, 0x30
14     pop rbp
15     ret
```


Stack Frame




```
1  auth:
2      ...
3      mov rdi, rax
4      call strncpy
5      mov byte [rbp-0x11], 0x00
6      ...
7
8  strncpy:
9      push rbp
10     mov rbp, rsp
11     sub rsp, 0x30
12     ...
13     add rsp, 0x30
14     pop rbp
15     ret
```

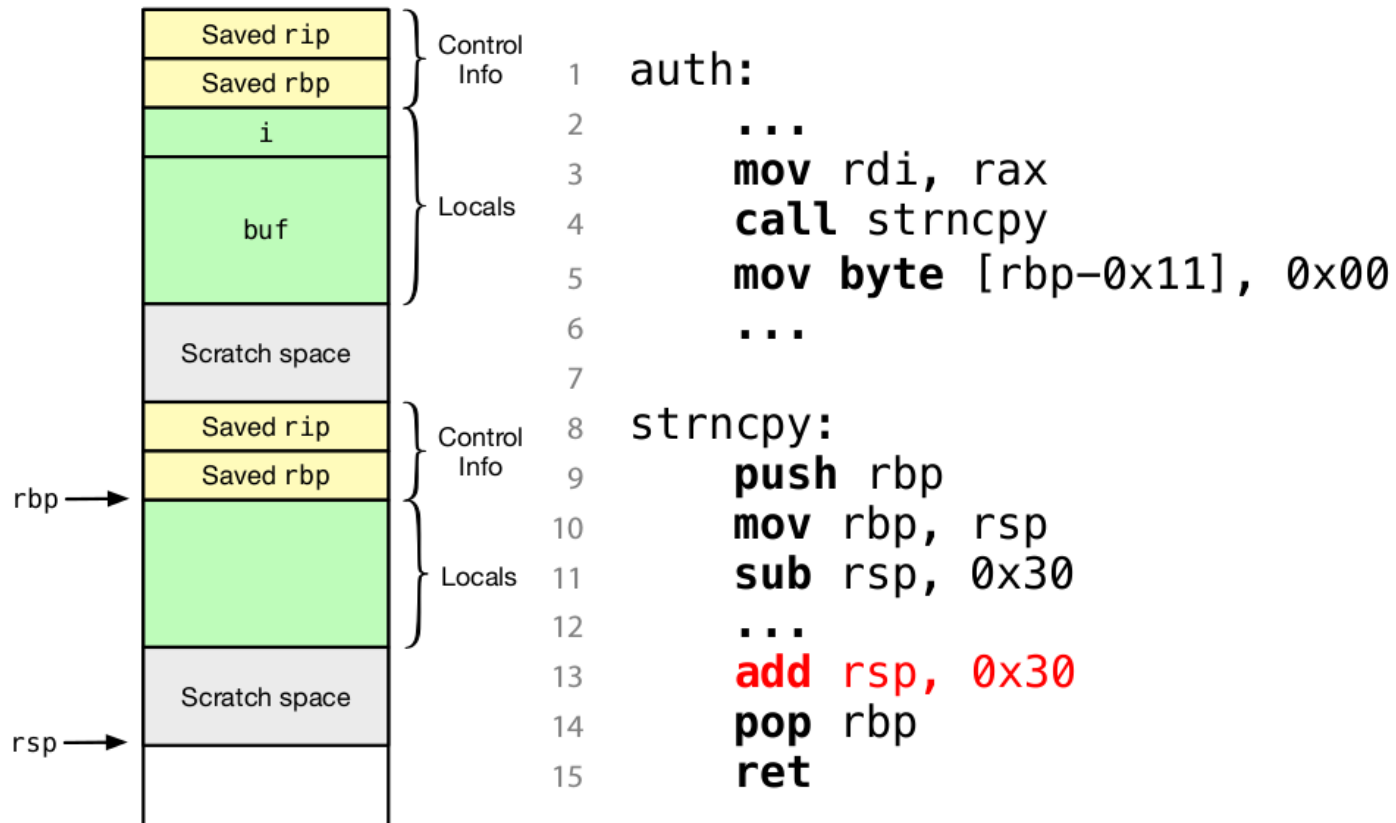
Stack Frame




Stack Frame




Stack Frame



Stack Frame




Stack Frame



```
1 auth:
2     ...
3     mov rdi, rax
4     call strncpy
5     mov byte [rbp-0x11], 0x00
6     ...
7
8 strncpy:
9     push rbp
10    mov rbp, rsp
11    sub rsp, 0x30
12    ...
13    add rsp, 0x30
14    pop rbp
15    ret
```

Stack Frame



```
1  auth:
2      ...
3      mov rdi, rax
4      call strncpy
5      mov byte [rbp-0x11], 0x00
6      ...
7
8  strncpy:
9      push rbp
10     mov rbp, rsp
11     sub rsp, 0x30
12     ...
13     add rsp, 0x30
14     pop rbp
15     ret
```

Procedure Arguments

Standards (*calling conventions*) exist for argument passing

- Specify where arguments are passed (registers, stack)
- Specify the caller and callee's responsibilities
 - Who deallocates argument space on the stack?
 - Which registers can be clobbered, and who must save them?

Why do we need standards?

- There are many ways to pass arguments
- How would code compiled by different developers and toolchains interoperate?

Calling Conventions

We often speak of *callers* and *callees*

- Caller: Code that invokes a procedure
- Callee: Procedure invoked by another function

Conventions must specify how registers must be dealt with

- Could always save them, but that is inefficient (why?)
- Usually, some registers can be overwritten (clobbered), others cannot
- Registers that can be clobbered: *caller* saved
- Registers that must not be clobbered: *callee* saved

SysV AMD64 ABI

x86_64 calling convention used on Linux, Solaris, FreeBSD, Mac OS X

- This is what you'll see most often in this course

First six arguments passed in registers

- rdi, rsi, rdx, rcx, r8, r9
 - Except syscalls, rcx → r10
- Additional arguments spill to stack

Return value in rax

cdecl

We've been concentrating on x86_64, but cdecl is important to know

- Linux 32 bit calling convention

Arguments

- Passed on the stack
- Pushed *right to left* (reverse order)

Registers

- eax, edx, ecx are *caller* saved
- Remainder are *callee* saved

Return value in eax

Caller deallocates arguments on stack after return

SysV AMD64 ABI Example

```
int auth( const char * user) {  
    size_t i;  
    char buf[16];  
    strncpy(buf, user, sizeof (buf));
```

auth:

push rbp	; save previous frame pointer
mov rbp, rsp	; set new frame pointer
sub rsp, 0x30	; allocate space for locals (i, buf)
movabs rdx, 0x10	; move sizeof(buf) to rdx
lea rax, [rbp-0x20]	; get the address of buf on the stack
mov qword [rbp-0x08], rdi	; move user pointer into stack
mov rsi, qword [rbp-0x08]	; move user pointer back into rsi
mov rdi, rax	; move buf into rdi
call strncpy	; call strncpy(rdi, rsi, rdx)
...	

Questions?