

# EC 440 – Introduction to Operating Systems

**Manuel Egele**

Department of Electrical &  
Computer Engineering  
Boston University

# Segmentation

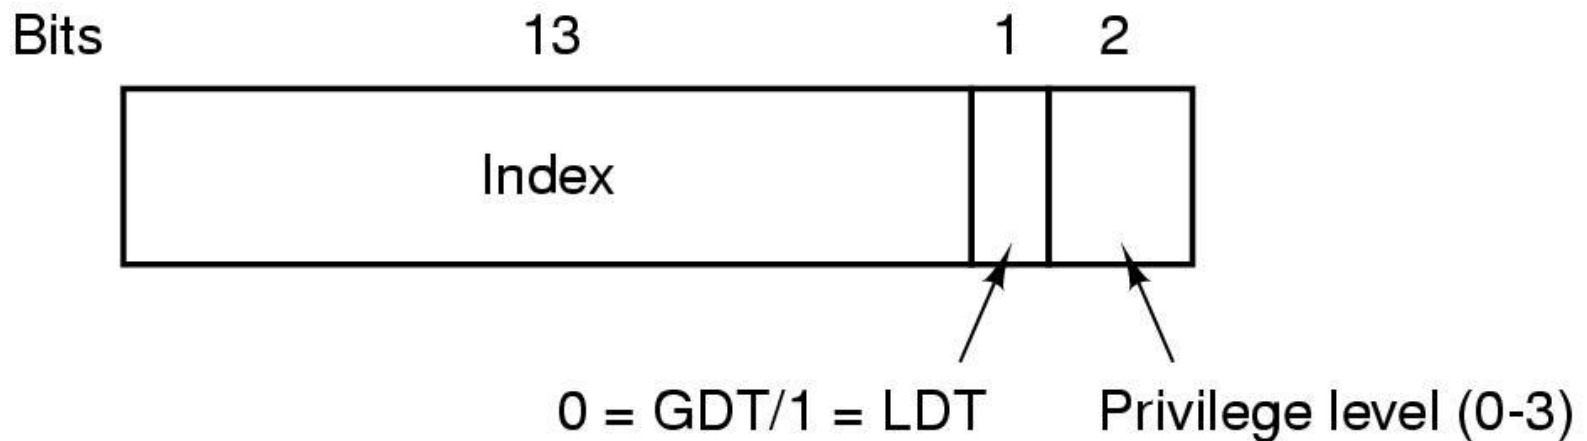
- One-dimensional address space is cumbersome to deal with if different portions of the program have to grow/shrink
- Provide the virtual machine with several independent address spaces, called *segments*
- Addressing is done by specifying
  - Segment
  - Address within the segment
- Advantages
  - Easy to share code and data segments (shared libraries)
  - Different segments can have different types of protection
- Segmentation is usually composed with paging

# Segmentation with Paging: Pentium

- Virtual memory with 16K segments
- Local Descriptor Table (LDT) for each program
- Global Descriptor Table (GDT) for the whole system
- Interrupt Descriptor Table (IDT) for the whole system
- To access a segment a selector for the segment is loaded into one of the segment registers (six in total)
  - CS holds code segment
  - DS holds data segment

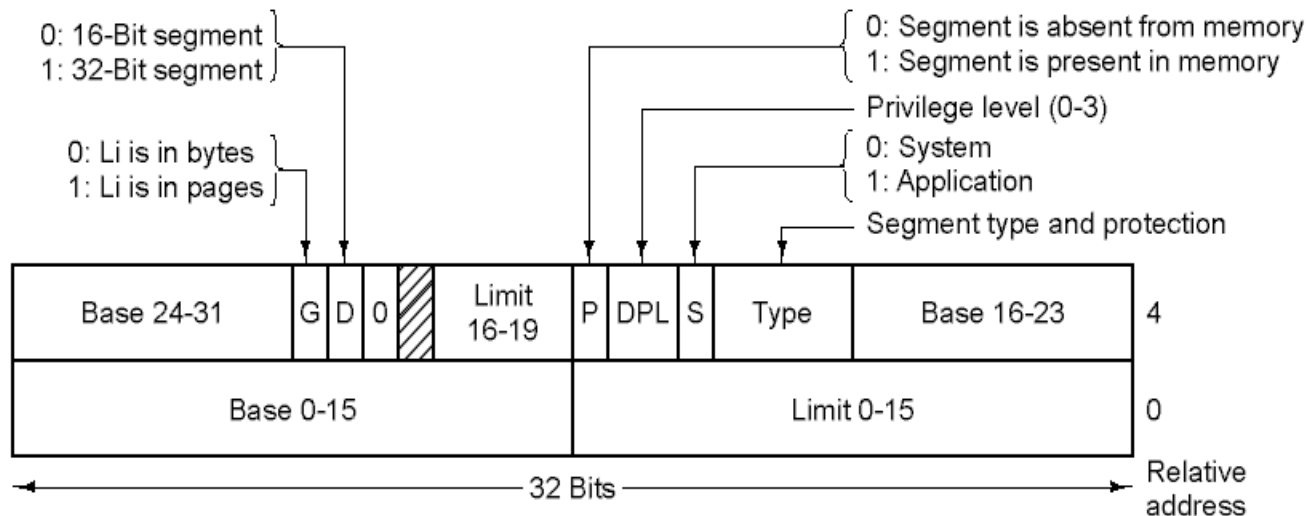
# Segment Selector

- A Pentium selector contains a bit to specify if the selector is part of the GDT or the LDT (8K segments each)
- A set of bits determines the privilege level
- Segment selector determines which segment descriptor to use



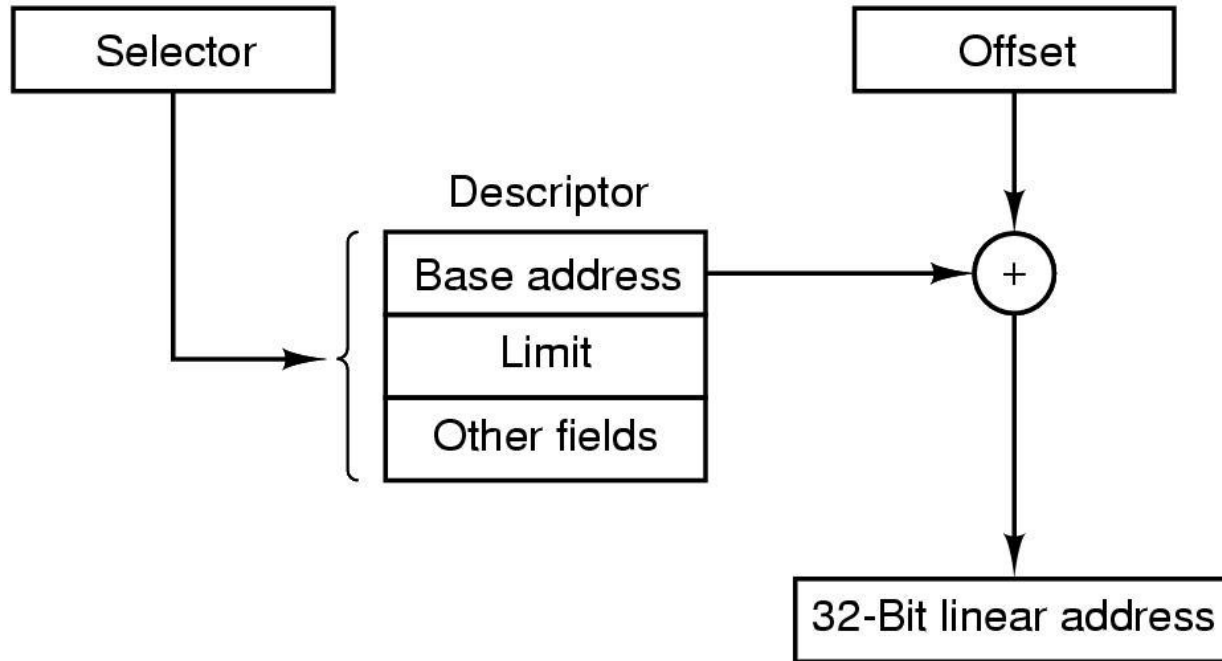
# Segment Descriptor

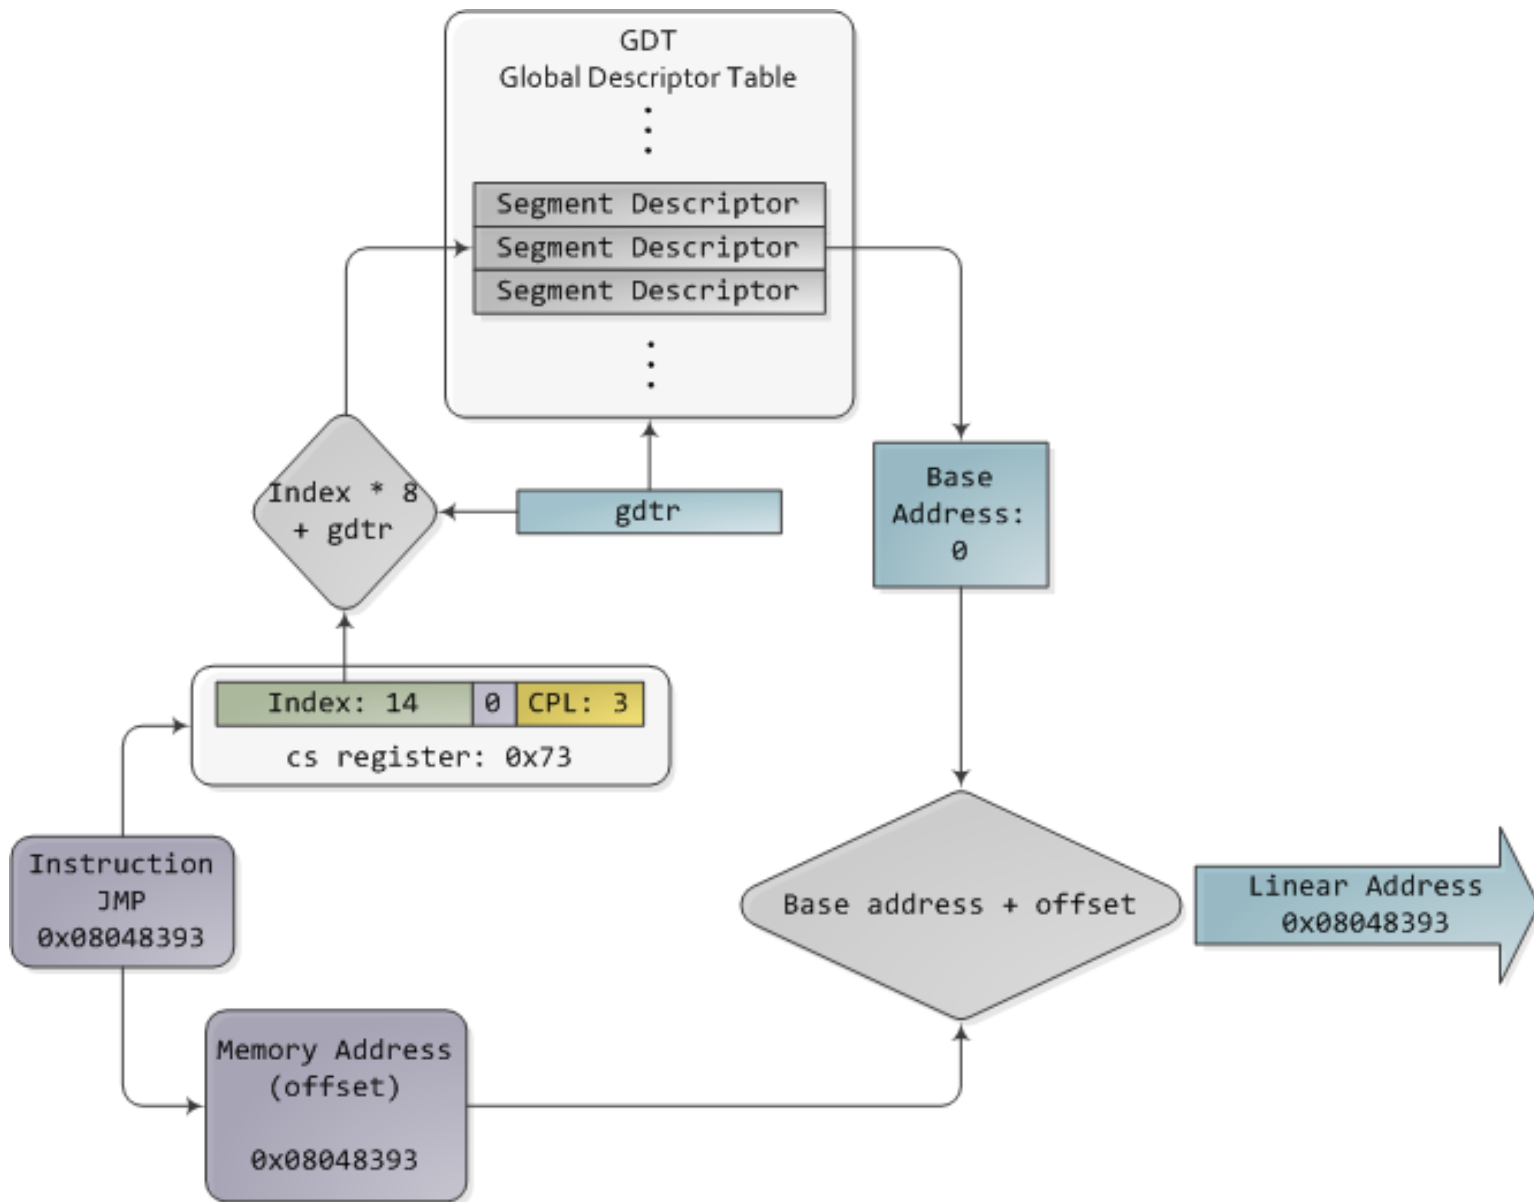
- The segment descriptor is 64 bit long
- The “limit” is expressed with 20 bits: if Granularity bit is 0, then max limit is 1MB; if the G-bit is 1, then limit is in pages of 4K (the missing 12 bits!)
- 3 Segment Types (code, data, system)



# Mapping An Address

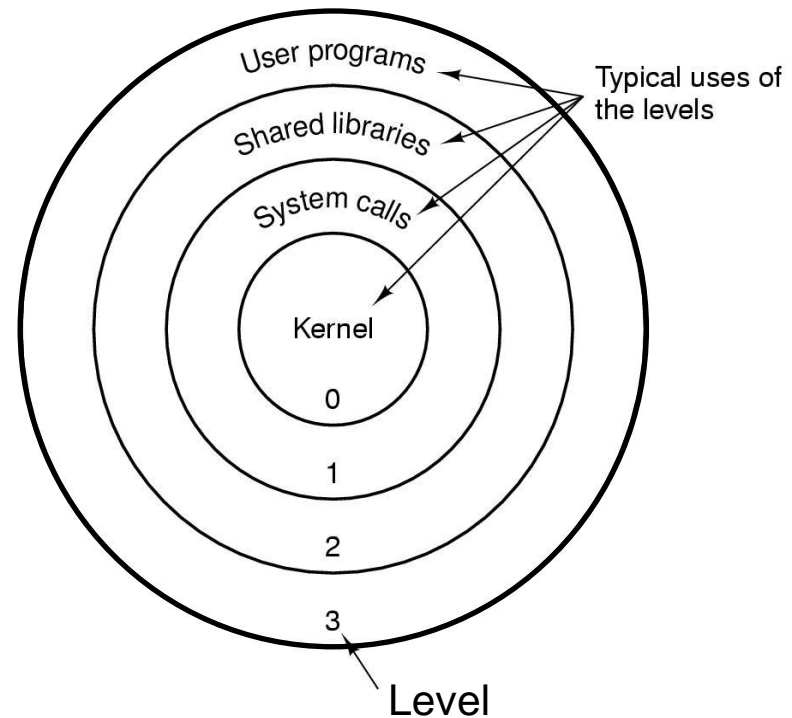
- Conversion of (selector, offset) pair to a linear address





# Protection on the Pentium

- Calls to procedures between protection levels must be performed by specifying a selector
- The selector is used to locate a *call gate* that gives the address of the required procedure
- This way, it is not possible to jump to arbitrary locations



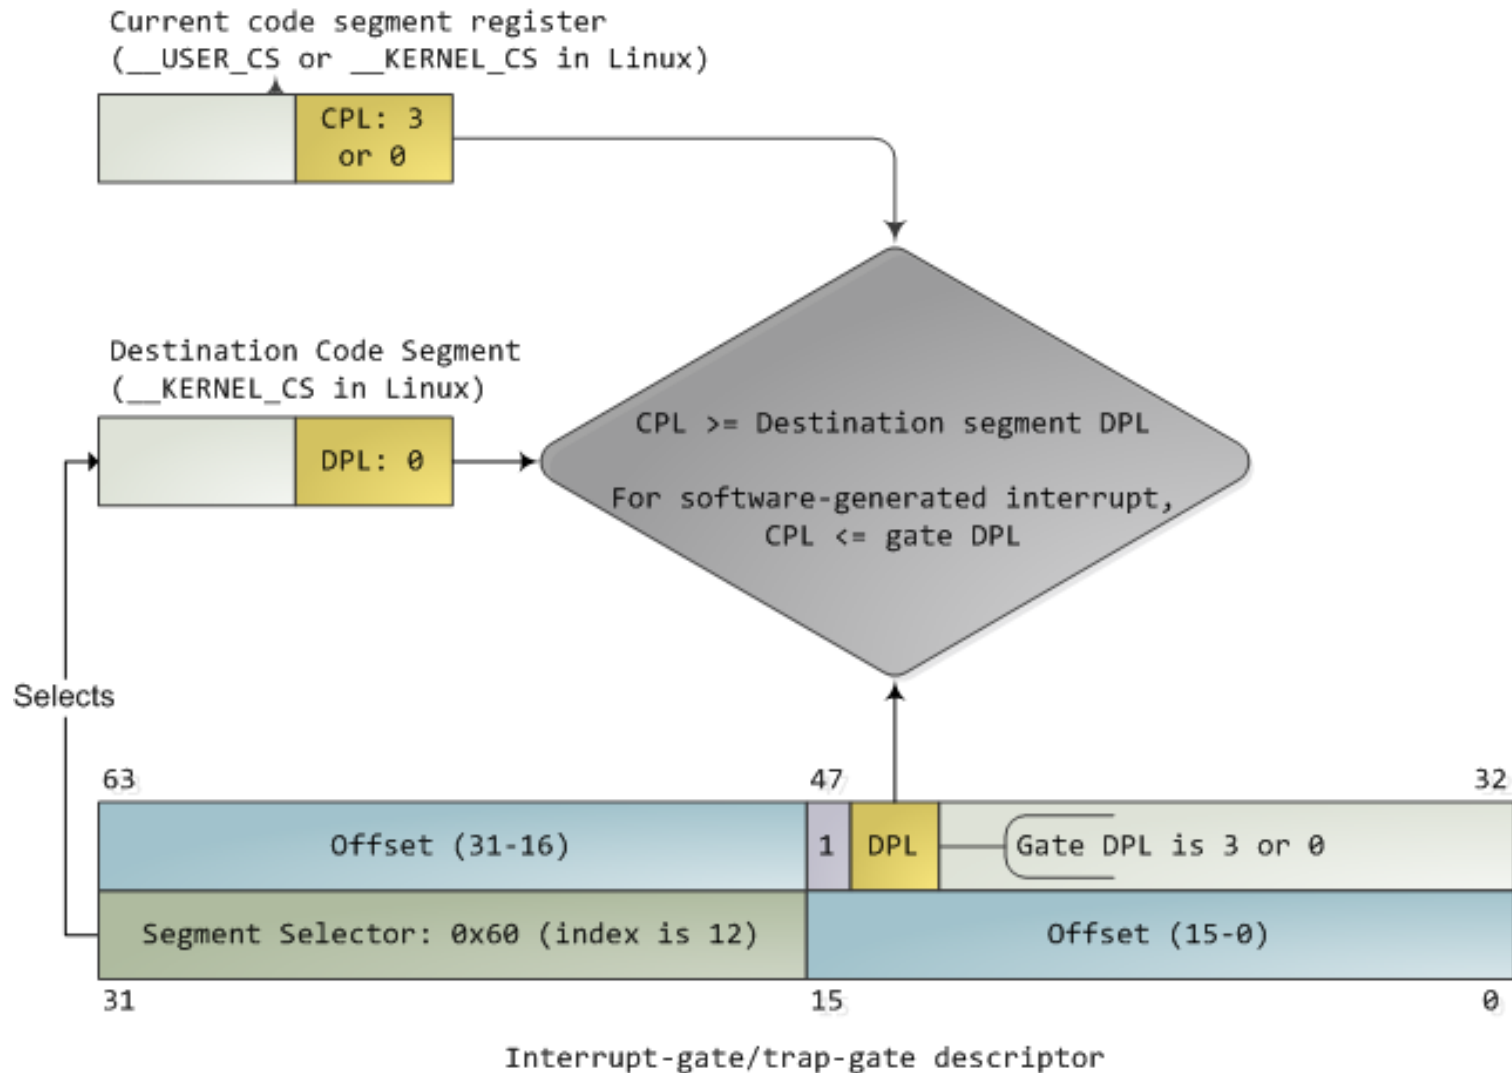


# Change Rings

- Jumping from user-space to kernel-space requires that the privilege level (i.e., the ring) is changed
- How does this work?
  - Segments!



# Segments & Interrupts / Syscalls



# INT 0x80 (system call) in Linux

```
/*
 * The default IDT entries which are set up in trap_init() before
 * cpu_init() is invoked. Interrupt stacks cannot be used at that point and
 * the traps which use them are reinitialized with IST after cpu_init() has
 * set up TSS.
 */
static const __initconst struct idt_data def_idts[] = {
    INTG(X86_TRAP_DE,      divide_error),
    ...

#ifdef CONFIG_IA32_EMULATION
    SYSG(IA32_SYSCALL_VECTOR, entry_INT80_compat),
#elif defined(CONFIG_X86_32)
    SYSG(IA32_SYSCALL_VECTOR, entry_INT80_32),
#endif
};
```

# INT 0x80 in Linux

```
struct idt_data {
    unsigned int    vector;
    unsigned int    segment;
    struct idt_bits bits;
    const void      *addr;
};

#define DPL0        0x0
#define DPL3        0x3

#define DEFAULT_STACK 0

#define G(_vector, _addr, _ist, _type, _dpl, _segment) \
{ \
    .vector      = _vector, \
    .bits.ist     = _ist, \
    .bits.type    = _type, \
    .bits.dpl     = _dpl, \
    .bits.p       = 1, \
    .addr         = _addr, \
    .segment      = _segment, \
}

/* System interrupt gate */
#define SYSG(_vector, _addr) \
    G(_vector, _addr, DEFAULT_STACK, GATE_INTERRUPT, \
      DPL3, __KERNEL_CS)
```

Can be called  
from ring 3

Loads a (kernel) code-  
segment with CPL = 0  
-> transitions to ring 0

**Questions?**